

The AI bottleneck isn't only chips. It's unused capacity.

An open marketplace thesis for bringing idle GPUs and TPUs online — at transparent cost, without overclaiming readiness.

Audience: infrastructure & AI platform leaders **Length:** ~4 min read

Contact: hellonecloudsmarketplace@googlegroups.com

Every board conversation about AI eventually lands on the same constraint: **compute**.

We talk about H100s and TPUs as if scarcity is absolute. In practice, scarcity is uneven. Hyperscalers are booked. Procurement cycles are slow. Quotes arrive late. Meanwhile, a quieter fact sits in plain sight:

A meaningful share of accelerators are already purchased — and idle.

They sit in neo-clouds, research labs, enterprises, and colocation rooms. Powered. Depreciating. Offline to the market that needs them most.

That gap — between *owned silicon* and *reachable capacity* — is the problem Neo Clouds is built to close.

The real failure mode

AI teams do not only lack chips. They lack **liquidity**.

Liquidity means:

- Can I find the right accelerator quickly?
- Can I trust it is actually online?
- Can I reserve it without a week of emails?
- Can I do that at a transparent listed rate?

Today, secondary capacity is fragmented. Buyers chase bilateral deals. Sellers have no clean distribution channel. The result is predictable: overpaying where capacity is

branded, and underusing capacity where it is not.

Spot markets inside a single cloud help utilization *inside that cloud*. They do not aggregate unused GPUs and TPUs across neo-clouds, labs, and independent operators into one discoverable catalog.

What Neo Clouds is

Neo Clouds is an open marketplace for unused accelerators.

The product thesis is deliberately simple:

1. **Unused** → **online** — bring idle GPUs and TPUs into a shared market
2. **Low / transparent cost** — listed rates buyers can compare
3. **Open** — open catalog, open APIs, open-source control plane

Customers browse by **hardware**, **workload intent** (training / inference / fine-tune), and **price** — not by opaque SKUs alone. Providers list capacity and, for live nodes, connect through a lightweight agent path: attestation, heartbeat, and provision acknowledgment.

That last point matters for executives evaluating seriousness. A listing on a webpage is not inventory. **Identity and liveness are different signals**. Attestation answers “what hardware claims to be here.” Heartbeat answers “is it online now.” Neo Clouds separates those on purpose.

What we are — and are not — claiming

Early access requires discipline.

Neo Clouds is live as a marketplace surface at neocloudsmarketplace.com. Demo inventory remains clearly labeled as simulated. Live provider nodes are a distinct path. **Payments are still disabled**. We would rather be precise than overclaim a paid fleet we have not finished wiring.

That honesty is strategy, not apology. Markets form when participants trust the rules. Overstating readiness destroys the only asset that matters in an early marketplace: credibility.

Who this is for

Providers with spare or intermittently idle GPUs/TPUs — neo-cloud operators, labs, enterprises with dark capacity, colo partners — who want another distribution channel without rebuilding a cloud.

Builders who need flexible accelerator access: short training bursts, fine-tunes, inference experiments — without treating every request as a hyperscaler procurement event.

If you sit between those worlds — infrastructure leadership, AI platform ownership, capacity planning — you already feel the inefficiency. Owned assets depreciating on one side. Roadmaps slipping on the other.

Why “open” is an executive decision

Closed marketplaces can work. Many do. But AI infrastructure is consolidating fast, and buyers are wary of new lock-in wearing a friendlier UI.

An open catalog and open APIs are not idealism. They are a bet that **secondary compute becomes infrastructure**, and infrastructure gets composed: schedulers, agents, internal portals, and future “sky” planners will want programmable supply — not another silo.

Neo Clouds is designed as a control plane for that supply: discover, reserve, connect. Orchestration depth, settlement, and hardened multi-tenancy come next. The wedge is liquidity.

The ask

If you run unused accelerators, or you are accountable for getting AI teams unblocked without burning budget on wait time:

Explore the marketplace: neocloudsmarketplace.com

Providers: start at the console / pilot path on the site

Contact: hellonecoloudsmarketplace@googlegroups.com

The next phase of AI infrastructure will not only be about who manufactures chips. It will be about who makes existing chips **reachable**.

Unused capacity that never reaches the market is a silent tax on every AI roadmap.

We are building the exchange that turns that tax into supply.

Neo Clouds · Executive LinkedIn article · Early access · Payments disabled

Source: [marketing/linkedin-executive-article.md](#)